

PH125.9x Capstone Project 1: MovieLens

Chris-FR

29 janvier 2019

Introduction

For this first project of the EDX PH125.9x capstone course, we will be creating a movie recommendation system using the MovieLens dataset. We will use the **10M records** version of this dataset. As a large number of students, one of the issue of this project was the performance of my laptop (2003, 4Gb RAM) to process the **edx** dataset.

In this report, we will:

- do a quick data exploratory analysis,
- then build 4 different models based on the **edx** dataset using Azure ML and the Penalized Least Squares regularization method
- and to finish we will compute the RMSEs of a **validation** dataset.

Data exploration and visualization

Loading data

To load the data, we used the code provided by the course page. We modified it to use the **fread** method of the **data.table** package.

The **edx** dataset has **9000055** rows and **6** columns.

The **validation** dataset has **999999** rows and **6** columns.

Movie genres dataset

We create a **movieGenres** matrix dataset containing the **movieId** as rows and genres as columns. We delete the “timestamp” and “title” columns from the **edx** and **validation** datasets.

```
# create the movieGenres dataset, containing the movieId as rows and the genres as columns
movieGenres <- edx %>% select(movieId, genres) %>%
  unique() %>%
  separate_rows(genres, sep = "\\|") %>%
  mutate(val=1) %>%
  spread(genres, val, fill=0)
# drop the "timestamp", "title" columns from edx and validation
edx$timestamp <- NULL
edx$title <- NULL
# we also drop the "genres" column from validation
validation$timestamp <- NULL
validation$title <- NULL
validation$genres <- NULL

head(movieGenres)
```

##	movieId (no genres listed)	Action	Adventure	Animation	Children	Comedy
## 1	1	0	0	1	1	1
## 2	2	0	0	1	0	0
## 3	3	0	0	0	0	1
## 4	4	0	0	0	0	1
## 5	5	0	0	0	0	1
## 6	6	0	1	0	0	0

```
##   Crime Documentary Drama Fantasy Film-Noir Horror IMAX Musical Mystery
## 1      0              0      0          1          0      0      0          0      0
## 2      0              0      0          1          0      0      0          0      0
## 3      0              0      0          0          0      0      0          0      0
## 4      0              0      1          0          0      0      0          0      0
## 5      0              0      0          0          0      0      0          0      0
## 6      1              0      0          0          0      0      0          0      0
##   Romance Sci-Fi Thriller War Western
## 1      0      0          0      0      0
## 2      0      0          0      0      0
## 3      1      0          0      0      0
## 4      1      0          0      0      0
## 5      0      0          0      0      0
## 6      0      0          1      0      0
```

To store the different RMSE results, we define a `rmse_results` tibble.

```
# creating the RMSE result table
rmse_results <- data_frame(method=character(0),
                           RMSE = numeric(0))
```

Data visualization

This section displays some of the graphs used to analyse the data.
We use the following consolidated data:


```
stats_users <- edx %>%
  group_by(userId) %>%
  summarize(countRating=n(), meanRating=mean(rating), medianRating=median(rating))

stats_movies <- edx %>%
  group_by(movieId) %>%
  summarize(countRating=n(), meanRating=mean(rating), medianRating=median(rating))

stats_rating <- edx %>%
  group_by(rating) %>%
  summarize(countRating=n())

stats_genres <- edx %>%
  separate_rows(genres, sep = "\\|") %>%
  group_by(genres) %>%
  summarize(countRating=n(), meanRating=mean(rating),
            medianRating=median(rating), countMovies=n_distinct(movieId))
```

Ratings distribution

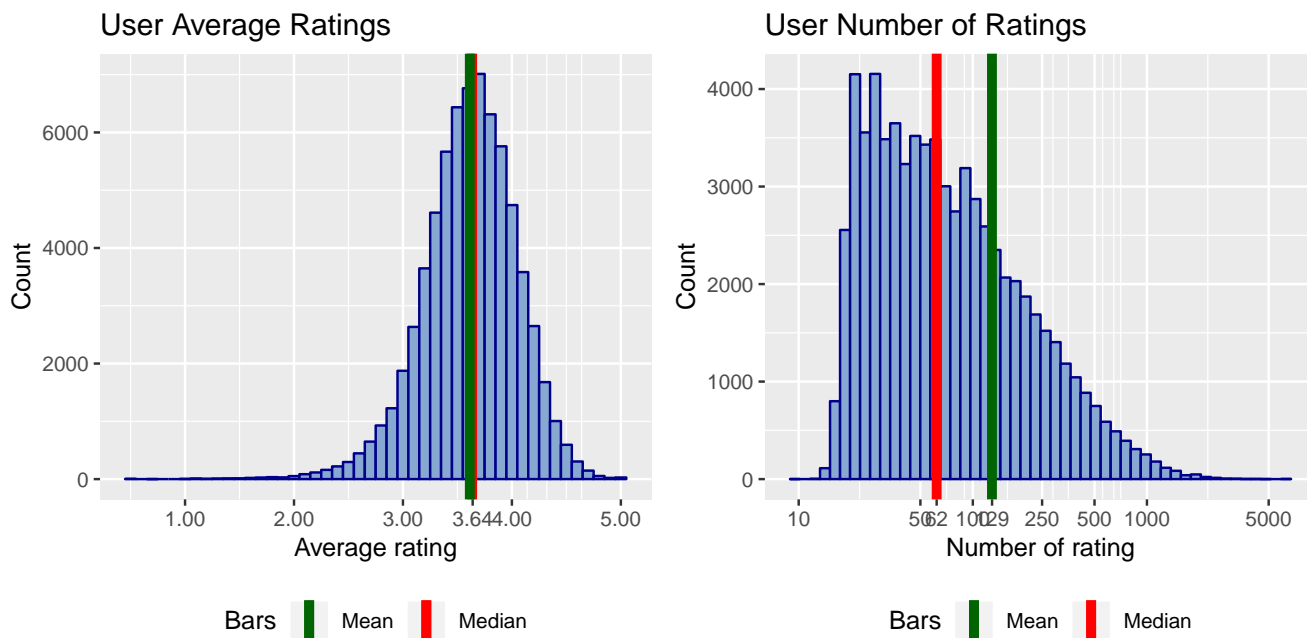


There are **9000055** ratings.

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	0.500	3.000	4.000	3.512	4.000	5.000

79.5% of these ratings are **whole** numbers. There is an unbalance between whole and half rating.

Users



There are **69878** users.

Number of ratings by users:

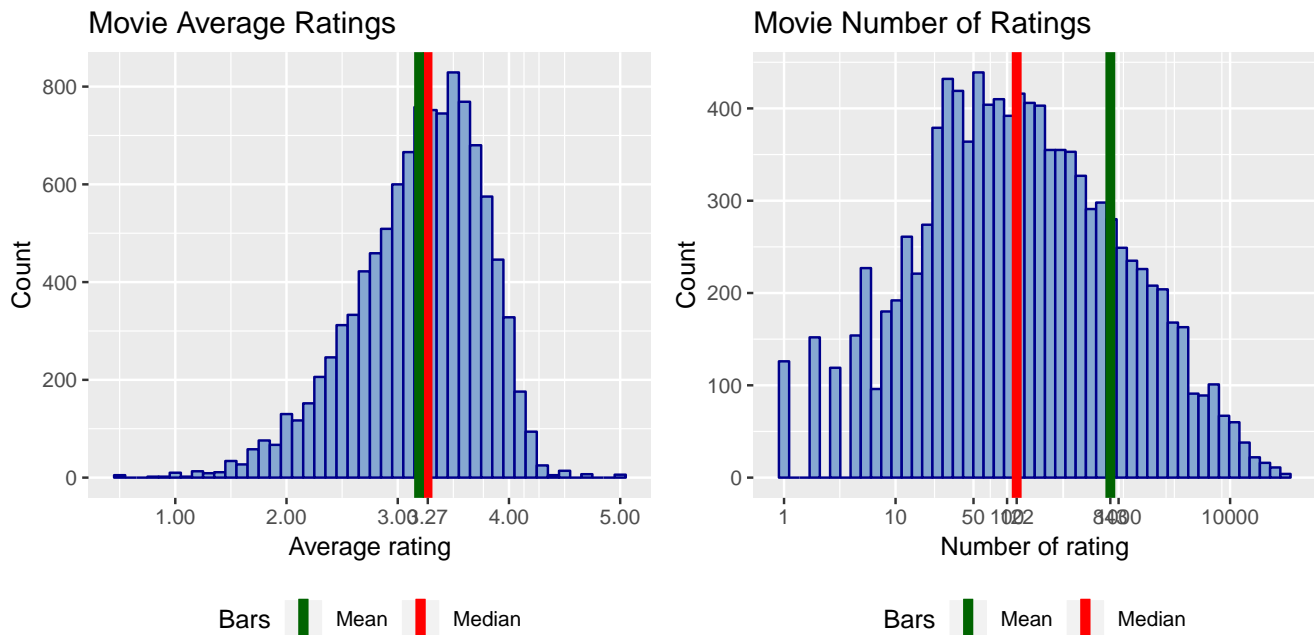
##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	10.0	32.0	62.0	128.8	141.0	6616.0

The mean number of rating by user is : **128.8**.

25% of the user have less than **32** ratings.

There are some very high value : **610** users rated more than 1000 movies, with a max number of rating of **6616**.

Movies



There are **10677** movies


Number of ratings by movies:

##	Min.	1st Qu.	Median	Mean	3rd Qu.	Max.
##	1.0	30.0	122.0	842.9	565.0	31362.0

The mean number of rating by movie is : **842.9**.

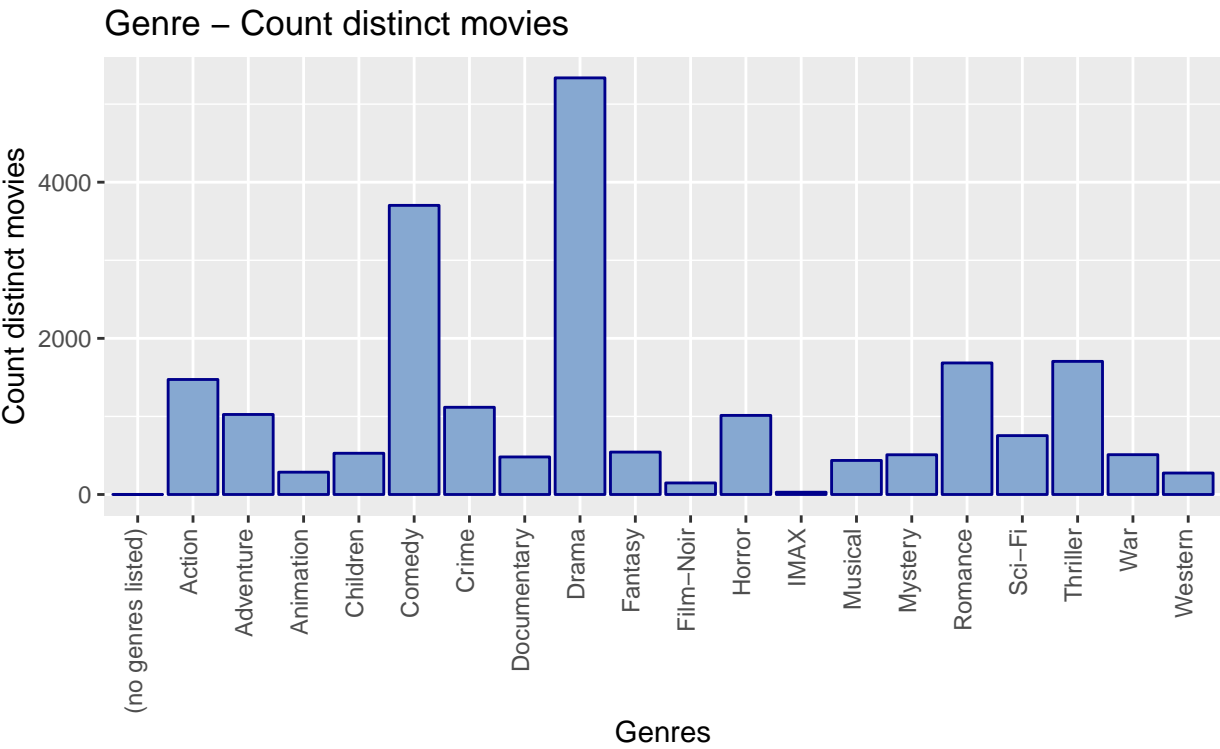
25% of the movies have less than **30** ratings. 25% of the movies have more than **565** ratings.

There are some very high value : **143** movies have more than 10000 movies, with a max number of rating of **31362**.

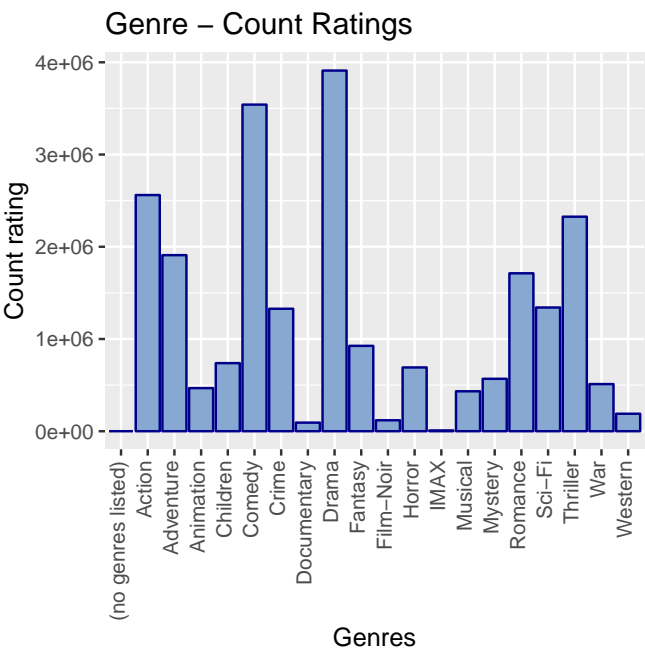



Movies with a high number of ratings seem to have higher means and less variabilities.

Genres



‘Drama’ and ‘Comedy’ are the most represented genres.



There is a genre effect.

Data analysis

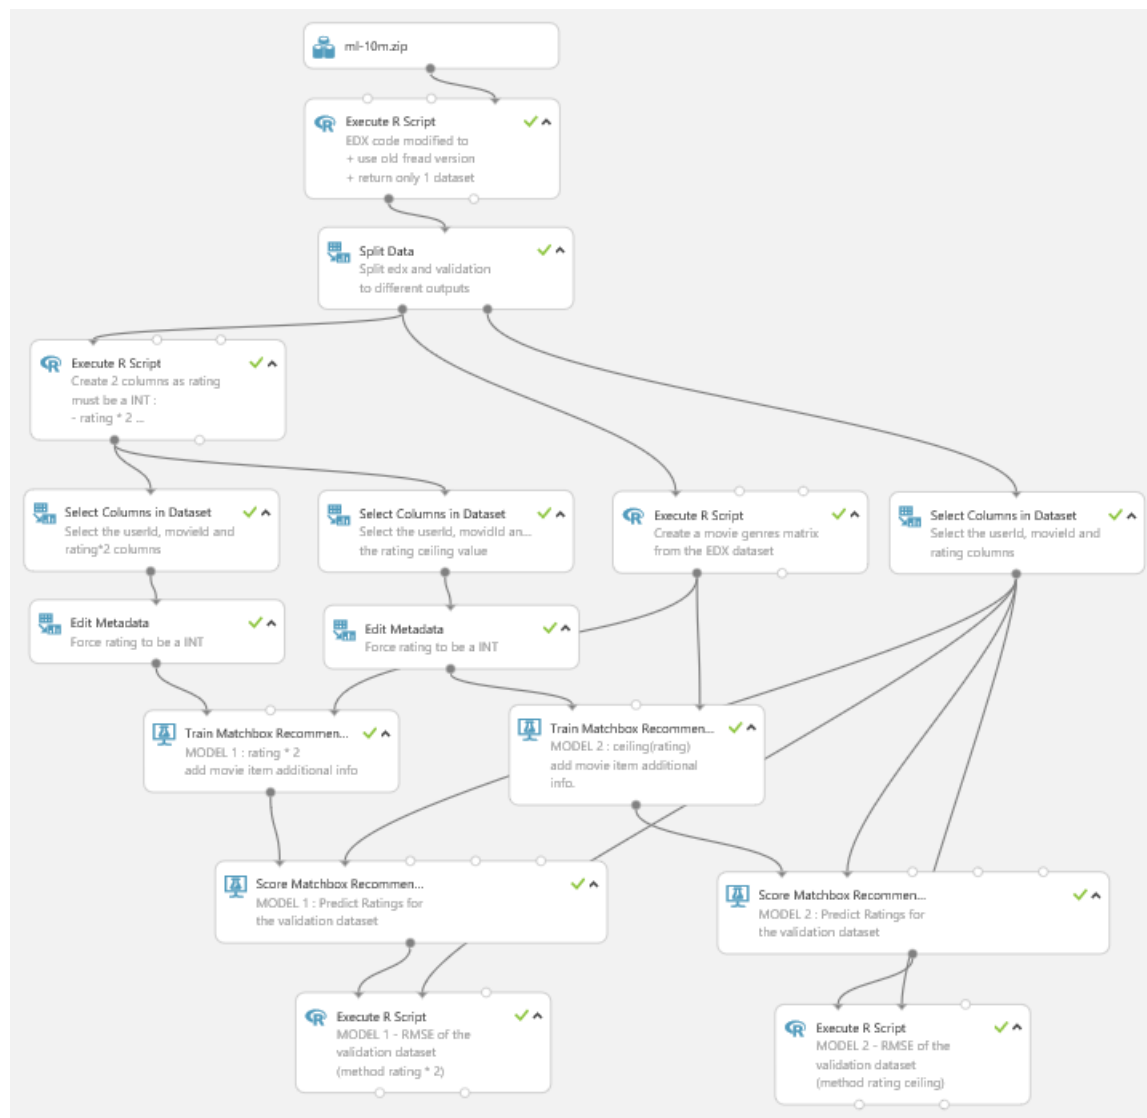
Even if it was out of this course scope, I found interesting to test Azure ML studio (with R integration).

Azure ML - Matchbox recommender

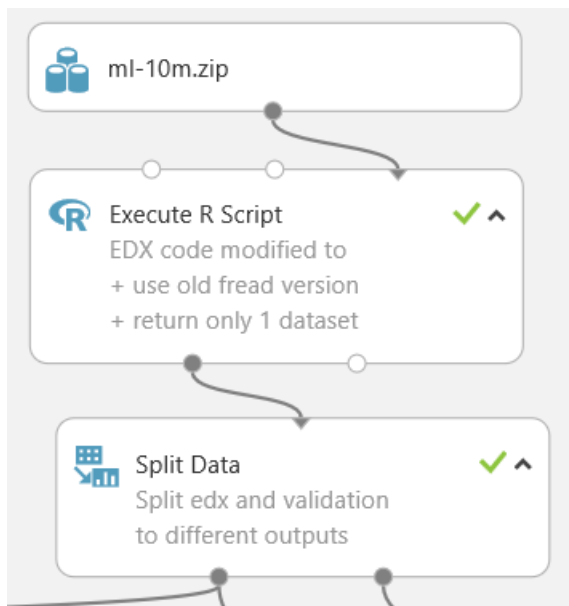
Microsoft provides a free trial access to test the Azure ML studio (100 modules by experiment, 10 Gb storage, 1 hour per experiment, single node execution).

Azure ML has a recommender module : the Matchbox Recommender.

Here is the schema of the full experiment I developed:



Loading the data



To load the data, an R script container containing the EDX code has been added.

3 main modifications have been done :

- the original ZIP file is not uploaded using HTTP. It was manually loaded as a dataset file. We used this dataset as the 3rd R script input. This input provides the unzipped file content in a SRC/ virtual directory.
- to speed up the **ratings** file loading, we used the **data.table fread** function. An issue was that in the available version of this package, the **text** parameter is unavailable and we have to separate the **readLines** function from the **fread** function.
- the R script module has only **one output**: to overcome this issue, I added a flag (train) and row binded the 2 datasets (edx and validation). We then used a **Split module** to split the **edx** and **validation** datasets using the **train** flag as split condition.

At the end of the R script, validation and edx datasets are concatenated, with an extra train column:

```

...
# use the old fread fonction : separate gsub and use paste.
ratings <- gsub("::", "\t", readLines("src/ml-10M100K/ratings.dat"))
ratings <- 'data.table'::fread(paste(ratings, collapse="\n"), col.names = c("userId", "movieId", "rating",
...
# add the train column and rbind the 2 datasets
edx$train<-1
validation$train<-0
data.set = rbind(edx, validation);
# Select data.frame to be sent to the output Dataset port
maml.mapOutputPort("data.set");

```



The left output of the **Split data** module is the **edx** dataset, the right output is the **validation** dataset.

Adding an item feature dataset












The 2nd and 3rd entries of the Matchbox recommender are a user features and item features datasets. For the item features dataset, we create a **genre** matrix dataset from the **edx** main dataset.

```
# Map 1-based optional input ports to variables
edx <- maml.mapInputPort(1) # class: data.frame
library(dplyr)
library(tidyr)
# create the genres matrix
movieGenres <- edx %>% select(movieId, genres) %>%
  unique() %>%
  tidyr::separate_rows(genres, sep = "\\|") %>%
  mutate(val=1) %>%
  spread(genres, val, fill=0)
# Select data.frame to be sent to the output Dataset port
maml.mapOutputPort("movieGenres");
```

EDX : MovieLens ratings ▶ Execute R Script ▶ Result Dataset

rows
10677

columns
21

	movieId	(no genres listed)	Action	Adventure	Animation	Children	Comedy	Crime	Documentary	Drama	Fantasy	Film- Noir
view as 												
	1	0	0	1	1	1	1	0	0	0	1	0
	2	0	0	1	0	1	0	0	0	0	1	0
	3	0	0	0	0	0	1	0	0	0	0	0

Train the Matchbox recommender

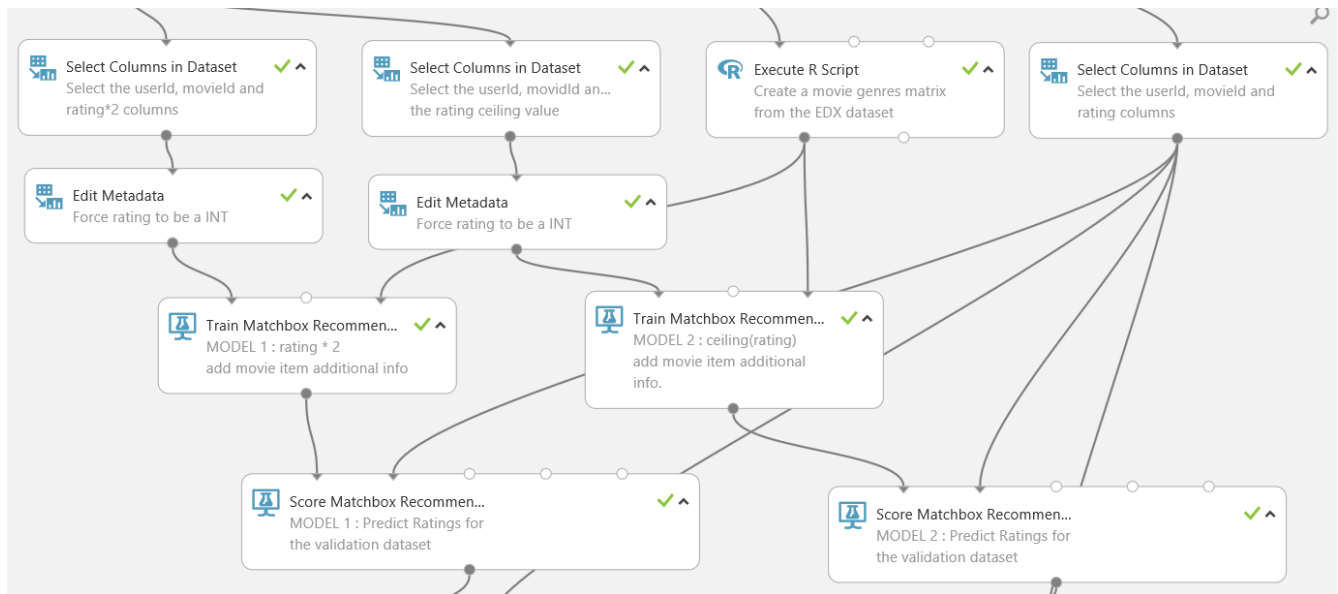
The ‘Train Matchbox Recommender’ module uses a dataset of user / item / rating entries. It returns a trained Matchbox recommender. We can then use the trained model in the Score Matchbox Recommender module to generate recommendations, find related users, find related items, compute expected ratings .

The rating of the user / item / rating dataset must be an **integer**.

I decided to test 2 cases (model 1 and 2):

- the ratings are multiplied by 2 (ratings from 1 to 10) - we keep the unbalance caused by the half notation numbers;
- the ratings are rounded using the R **ceiling** function.

```
# Map 1-based optional input ports to variables
dataset1 <- maml.mapInputPort(1) # class: data.frame
# convert rating
dataset1$ratingbytwo <- dataset1$rating * 2
dataset1$ratingceiling <- ceiling(dataset1$rating)
# Select data.frame to be sent to the output Dataset port
maml.mapOutputPort("dataset1");
```

The train module parameters were set to default.

what was not done :

In this version of the experiment, the full edx dataset was used to train the models. I used split validation on another experiment but did not on this one.

I also did not try to tune the recommender parameters (the parameters are the default ones).

Regularized Movie + User Effect Model

The 3rd solution tested in this report is the **Regularized Movie + User Effect Model**.

Finding lambdas with 5-folds CV

We will compute 2 different lambdas as the penalties may not be the same between the movie and user effects :

- li : for the movie effect
- lu : for the user effect

To pick the best values of these 2 lambdas, we will do a **5-folds cross validation**. We will split edx is 5 folds, then take 4 folds as train set and 1 fold as test set and process every combination. We will then take the values having the better RMSE result average accross the 5 folds.

To avoid the cases where the test set contains a movie or a user not present in the train set, we will use the same code than the first edx / validation creation (using a temp dataset first).

```

# create a dataframe to store each result
df.result <- data.frame( "k" = integer(0),
                        "li" = numeric(0),
                        "lu" = numeric(0),
                        "RMSE" = numeric(0))

# we are doing k-folds cross validation
k<-5

# create the k folds with the CARET package
set.seed(1)
folds <- createFolds(edx$rating, k=k, returnTrain = FALSE)

# for the report compute only 3 to 7
lambdasI <- seq(3, 7, 0.25)

```

```

lambdasU <- seq(3, 7, 0.25)

# for each of the 10 folds:
for (i in 1:k){

  # print(paste("Starting fold:" , i))

  # separate the train / test set using the fold #i
  train_set <- edx[-folds[[i]],]
  temp <- edx[folds[[i]],]

  # Make sure userId and movieId in the test set are also in train set
  test_set <- temp %>%
    semi_join(train_set, by = "movieId") %>%
    semi_join(train_set, by = "userId")

  # Add rows removed from test_set set back into train_set set
  removed <- anti_join(temp, test_set)
  train_set <- rbind(train_set, removed)
  rm(removed, temp)

  print(paste("Computing rmse for fold:" , i, dim(train_set)[1], dim(test_set)[1]))

  mu <- mean(train_set$rating)

  b_i <- train_set %>%
    group_by(movieId) %>%
    summarize(sum_b_i = sum(rating - mu), n=n())

  for(li in lambdasI){

    b_i$b_i <- b_i$sum_b_i/(b_i$n+li)

    b_u <- train_set %>%
      left_join(b_i, by="movieId") %>%
      group_by(userId) %>%
      summarize(sum_b_u = sum(rating - b_i - mu), n=n())

    for(lu in lambdasU){

      b_u$b_u <- b_u$sum_b_u / (b_u$n+lu)

      predicted_ratings <-
        test_set %>%
        left_join(b_i, by = "movieId") %>%
        left_join(b_u, by = "userId") %>%
        mutate(pred = mu + b_i + b_u) %>%
        .$pred


      myrmse <- RMSE(predicted_ratings, test_set$rating, na.rm = T)

      # print(paste(li, lu, myrmse))
      df.result[nrow(df.result) + 1,] = list(k=i, li=li, lu=lu, RMSE=myrmse)
    }
  }
}

```

```
## [1] "Computing rmse for fold: 1 7200090 1799965"
## [1] "Computing rmse for fold: 2 7200087 1799968"
## [1] "Computing rmse for fold: 3 7200080 1799975"
## [1] "Computing rmse for fold: 4 7200087 1799968"
## [1] "Computing rmse for fold: 5 7200083 1799972"
```

li	lu	minRMSE	meanRMSE	maxRMSE
4.5	5	0.8646528	0.8654837	0.8662331



Mean,

Min and Max of 3 li and theirs lu over the 5-folds.

Compute the full model values (μ , b_i , b_u) with the best li and lu lambdas

Computing the model with li = 4.5 and lu = 5.

```
mu <- mean(edx$rating)

b_i <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+li))

b_u <- edx %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lu))
```

Regularized Movie + User Effect + Genre effect Model

The last solution tested in this report is the Regularized Movie + User Effect + Genre Effect Model.

As the computation of the optimal lambdas is very slow on my laptop (+5 hours), we will use the previously calculated lambdas (λ_u and λ_i) and just add a genre effect to the result (i.e. as if we start from the residuals of the last models).

The genre effect is given by the following formula:

$$Y_{u,i} = \mu + b_i + b_u + \sum_{k=1}^K x_{u,i}^k \beta_k + \varepsilon_{u,i}, \text{ with } x_{u,i}^k = 1 \text{ if } g_{u,i} \text{ is genre } k.$$

To simplify the model, we will group the movies using a **genre clustering**. A movie is now part on an unique cluster group (and not to multiple genres). The formula becomes:

$$Y_{u,i} = \mu + b_i + b_u + b_{u,g} + \varepsilon_{u,i}$$

For the penalty, we will use the same than the user one.

Genres cluster

```
colSums(movieGenres[2:21])
```

## (no genres listed)	Action	Adventure
## 1	1473	1025
## Animation	Children	Comedy
## 286	528	3703
## Crime	Documentary	Drama
## 1117	481	5336
## Fantasy	Film-Noir	Horror
## 543	148	1013
## IMAX	Musical	Mystery
## 29	436	509
## Romance	Sci-Fi	Thriller
## 1685	754	1705
## War	Western	
## 510	275	

As there is only one record of the “(no genres listed)”, we remove this value.

We then compute the euclidean distance matrix and build out a Hierarchical Cluster.


```
movieGenres$`(no genres listed)` <- NULL

# compute the distance
# no need to normalize the values are they have the same scale
distances <- dist(movieGenres[2:20], method="euclidean")
# Hierarchical Cluster
clusterMovies <- hclust(distances, method="ward.D")
```

We split in 10 groups.

```
plot(clusterMovies, labels=FALSE)
clusterGroups <- cutree(clusterMovies, k = 10)
rect.hclust(clusterMovies, k=10, border="red")
```

Cluster Dendrogram



distances
hclust (*, "ward.D")

Here is a view of the mean of each genres in each cluster:

```
view_cluster_means <- sapply(1:10, function(x){round(colMeans(movieGenres[clusterGroups==x,2:20]),2)})
view_cluster_means #>% knitr::kable()
```

	[,1]	[,2]	[,3]	[,4]	[,5]	[,6]	[,7]	[,8]	[,9]	[,10]
## Action	0.32	0.0	0	0.09	0	0.00	0	0.00	0.28	0
## Adventure	0.24	0.0	0	0.02	0	0.00	0	0.01	0.16	0
## Animation	0.07	0.0	0	0.00	0	0.00	0	0.00	0.00	0
## Children	0.13	0.0	0	0.00	0	0.00	0	0.00	0.00	0
## Comedy	0.33	1.0	1	0.10	0	0.00	0	0.06	0.11	1
## Crime	0.21	0.0	0	0.04	0	0.46	0	0.01	0.01	0
## Documentary	0.01	0.0	0	0.00	0	0.00	0	1.00	0.03	0
## Drama	0.35	0.4	0	0.11	1	0.79	1	0.05	0.76	1
## Fantasy	0.14	0.0	0	0.00	0	0.00	0	0.00	0.01	0
## Film-Noir	0.04	0.0	0	0.00	0	0.00	0	0.00	0.00	0
## Horror	0.06	0.0	0	1.00	0	0.00	0	0.02	0.01	0
## IMAX	0.00	0.0	0	0.00	0	0.00	0	0.04	0.00	0
## Musical	0.11	0.0	0	0.00	0	0.00	0	0.00	0.01	0
## Mystery	0.11	0.0	0	0.10	0	0.00	0	0.00	0.01	0
## Romance	0.15	1.0	0	0.00	0	0.00	1	0.01	0.12	0
## Sci-Fi	0.16	0.0	0	0.17	0	0.00	0	0.00	0.02	0
## Thriller	0.24	0.0	0	0.36	0	0.74	0	0.00	0.07	0
## War	0.01	0.0	0	0.00	0	0.00	0	0.01	0.95	0
## Western	0.07	0.0	0	0.00	0	0.00	0	0.00	0.00	0

We save the pair : movieId / cluster.

Then we add the **gcluster** (genre cluster) to the **edx** and **validation** dataset. (this step will save us a JOIN each time).

```

movieGenres$gcluster <- clusterGroups
movieGenres <- movieGenres %>% select(movieId, gcluster)
# clean up
rm(distances, clusterMovies, view_cluster_means, clusterGroups)
# add the cluster group to EDX and validation dataset
edx <- edx %>%
  inner_join(movieGenres, by = "movieId")
edx$genres <- NULL
head(edx)

```

```

##   userId movieId rating gcluster
## 1      1      122      5         2
## 2      1      185      5         1
## 3      1      292      5         1
## 4      1      316      5         1
## 5      1      329      5         1
## 6      1      355      5         1

```

```

validation <- validation %>%
  inner_join(movieGenres, by = "movieId")
head(validation)

```

```

##   userId movieId rating gcluster
## 1      1      231      5         3
## 2      1      480      5         1
## 3      1      586      5         1
## 4      2      151      3         1
## 5      2      858      2         6
## 6      2     1544      3         4

```

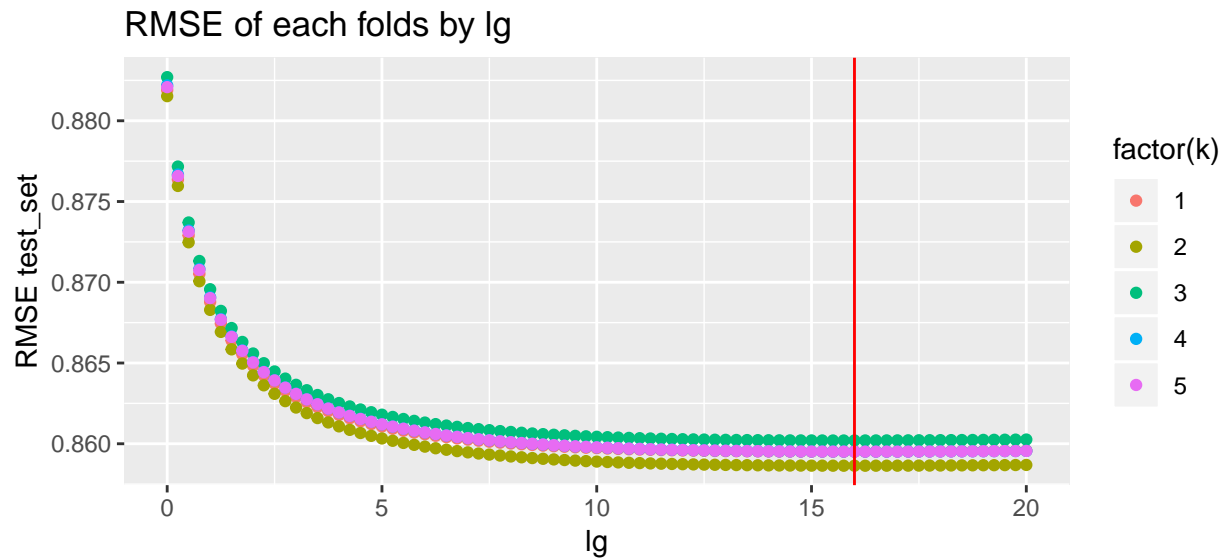
Genres effect

We can now compute the genres effect on the `edx` dataset.

As stated in the introduction of this section, we will use the residuals of the previous model (we use `li` and `lu` values of the model 3 and do not recompute them) and we will use 5-folds cross validation to compute the penalization lambda of the genre (`lg`).

The best value for `lg` is:

<code>lg</code>	minRMSE	meanRMSE	maxRMSE
16	0.8586403	0.8594774	0.8602032



Compute the full model values (μ , b_i , b_u , b_g)

We can now compute the full model on the edx dataset:

```
print(paste('li=',li, 'lu=', lu, 'lg=' , lg))

## [1] "li= 4.5 lu= 5 lg= 16"

mu <- mean(edx$rating)

b_i <- edx %>%
  group_by(movieId) %>%
  summarize(b_i = sum(rating - mu)/(n()+li))

b_u <- edx %>%
  left_join(b_i, by="movieId") %>%
  group_by(userId) %>%
  summarize(b_u = sum(rating - b_i - mu)/(n()+lu))

b_g <- edx %>%
  left_join(b_i, by="movieId") %>%
  left_join(b_u, by = "userId") %>%
  group_by(userId, gcluster) %>%
  summarize(b_g = sum(rating - b_i -b_u - mu)/(n()+lg))
```

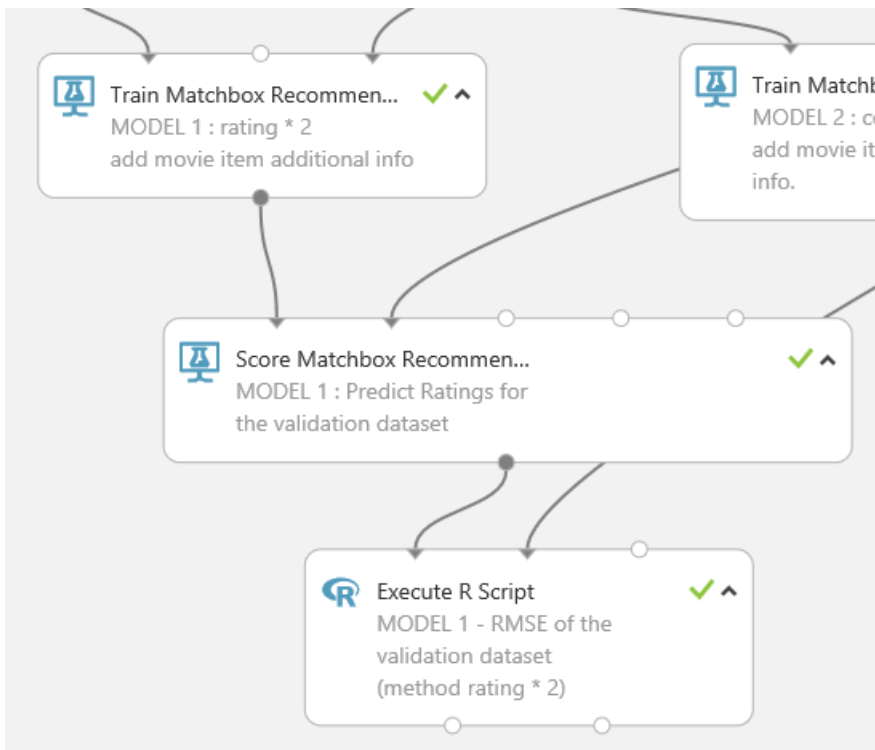
Results

We can now try the finalized models on our validation dataset.

Azure ML (model 1 and 2)

To compute the final RMSE, we use the validation dataset and compute its predicted ratings using the Score Matchbox module.

The final RMSE is computed in the last R script module.



The RMSE is available in the output dataset.

```

# Map 1-based optional input ports to variables
pred <- maml.mapInputPort(1) # class: data.frame
val <- maml.mapInputPort(2) # class: data.frame

error <- val$rating - pred$Rating
myrmse <- sqrt(mean(error^2))

# Select data.frame to be sent to the output Dataset port
myreturn<- data.frame(RMSE=myrmse)
maml.mapOutputPort("myreturn");

```

RMSE

0.939061

method	RMSE
Azure ML Matchbox - ratings*2	1.015422
Azure ML Matchbox - ceiling	0.939061

Regularized Movie + User Effect Model (model 3)

```

predicted_ratings <-
  validation %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  mutate(pred = mu + b_i + b_u ) %>%

```



```

    .$pred

rmse_model3 <- RMSE(predicted_ratings, validation$rating)

rmse_results <- dplyr::bind_rows(rmse_results,
                                data_frame(method="Regularized Movie + User Effect Model",
                                            RMSE = rmse_model3))

rmse_model3

## [1] 0.8648192

```

Regularized Movie + User Effect + Genre effect Model (model 4)

For the genre effect, we just need to ensure that if the user / genrecluster does not exist, the value of the average `b_g` is 0.

```

predicted_ratings <-
  validation %>%
  left_join(b_i, by = "movieId") %>%
  left_join(b_u, by = "userId") %>%
  left_join(b_g, by = c("userId", "gcluster")) %>%
  mutate(b_g = replace_na(b_g, 0)) %>%
  mutate(pred = mu + b_i + b_u + b_g) %>%
  .$pred

rmse_model4 <- RMSE(predicted_ratings, validation$rating)

rmse_results <- dplyr::bind_rows(rmse_results,
                                data_frame(method="Regularized Movie + User Effect + Genre effect Model",
                                            RMSE = rmse_model4))

rmse_model4

## [1] 0.8579097

```

Conclusion

The 4 RMSEs for the `validation` dataset are the following:

method	RMSE
Azure ML Matchbox - ratings*2	1.0154220
Azure ML Matchbox - ceiling	0.9390610
Regularized Movie + User Effect Model	0.8648192
Regularized Movie + User Effect + Genre effect Model	0.8579097

As expected from the average of the cross validations, the model with the best result is the ‘Regularized Movie + User Effect + Genre effect Model’ with an RMSE of **0.8579**. This value is near the ones computed using 5-folds cross validation on the `edx` dataset: average = 0.8595 (min = 0.8586, max = 0.8602).

I wish i had a better hardware to test more methods (SVM,...).